

Phongsakon Mark Konrad

+45 55 24 01 58 phongsakon@outlook.dk phomarkon.github.io
linkedin.com/in/phongsakonmarkkonrad github.com/phomarkon Google Scholar
Sønderborg, Denmark

Summary

I study how and why models learn, how what they learn is stored and represented, and how we get to models we can actually trust, with mechanistic interpretability as one lens among several. I run the full research lifecycle myself: problem formulation, experimental design, pipeline engineering, manuscript writing, and venue selection. Current work spans circuit-level audits of routing behaviour installed by fine-tuning (under review), evidential standards for safety claims, causal tracing of moral reasoning, and a cross-family study of chain-of-thought safety, alongside work on vision, medical-imaging, and remote-sensing models. The questions, not any one model family, are what I am after.

My conviction is simple: we do not need smarter models, we need more aligned ones. My goal from here is a PhD. I do not yet know which of my questions is the one I will spend years on, and I am comfortable with that for now, because I am following the curiosity and the questions keep getting sharper. I am a final-year BSc Software Engineering student at the University of Southern Denmark, heading to the MPhil in Machine Learning and Machine Intelligence at the University of Cambridge in 2026, and I bring an unconventional background (four years of military service, three startups, and a semester at HKUST) that shapes how I identify problems worth solving.

Education

University of Cambridge Starting Oct 2026
MPhil in Machine Learning and Machine Intelligence Cambridge, United Kingdom

The Hong Kong University of Science and Technology (HKUST) Aug 2025 – Dec 2025
Exchange Semester, Dept. of Computer Science and Engineering Hong Kong SAR, China

- Specialisation courses: COMP4211 Machine Learning, COMP4471 Deep Learning in Computer Vision, COMP4901B Large Language Models, COMP4901Z Reinforcement Learning.
- Postgraduate course: COMP6411D Data Visualisation.

University of Southern Denmark (SDU) Sep 2023 – Jun 2026
Bachelor of Science (BSc), Software Engineering Sønderborg, Denmark

- Expected graduation: June 2026. Current GPA: $\approx 10.6/12$ ($\approx 3.7/4.0$ US GPA).
- Practical, project-centric curriculum with mandatory semester-long team projects on complex, data-intensive software systems in domains such as IoT and AI, often in collaboration with industry partners.
- BSc thesis (in progress, with T. L. Adam): *The Trader's Trinity: Forecasting Models, RL Agents, and LLM Judges for Day-Ahead Markets*.

Selected Publications

2026 **P. M. Konrad**, T. Tanel, and S. Ayvaz. Self-Reports Do Not Identify Self-Models: An Identifiability Test for Counterfactual Reports. Accepted (oral), *PhilML Workshop, ICML 2026*. [PDF]
TL;DR. Self-reports are evidence about behavior in a prompt, not proof of a self-model. Wrong-source demonstrations pull affect-state reports toward the source family unless mechanism binding holds.

2026 **P. M. Konrad**, T. Tanel, and S. Ayvaz. Decoded but Unused: Instruction Tuning Routes Moral Framing into the Judgment Readout. Accepted (poster), *ICML 2026 Workshop on Mechanistic Interpretability*. [PDF]
TL;DR. Moral framing is linearly decodable in pretrained models but only causally routed to judgment after instruction tuning.

- 2026 **P. M. Konrad**, T. Tanel, and S. Ayvaz. A Path Already Walked: On Inheriting Network-Neuroscience Tools for Mechanistic Interpretability. Accepted (virtual poster), *ICML 2026 Workshop on Mechanistic Interpretability*. [PDF]
TL;DR. Position: mech-interp should import the graph vocabulary of network neuroscience (modularity, rich clubs, motifs) under explicit contracts.
- 2026 **P. M. Konrad**, T. Tanel, and S. Ayvaz. Routing Subspaces: Auditing Evaluation-to-Deployment Mismatch in Fine-Tuned Language Models. Under review. [PDF]
TL;DR. Path-patching localises an eval-vs-deployment signal to a narrow mid-depth attention window. Clamping a low-dimensional subspace closes the gap in 11 of 12 cells.
- 2026 **P. M. Konrad**, T. Tanel, and S. Ayvaz. Acceptance Cards: A Four-Diagnostic Standard for Safe Fine-Tuning Defense Claims. Under review. arXiv:2605.10575
TL;DR. A four-gate evidential standard for safe-fine-tuning defense claims. SafeLoRA fails the full-card pass under matched-action audit.

Full Publication List

P. M. Konrad is the corresponding first author except where noted. Workshop papers are listed separately from work under review; nothing below is double-counted.

Peer-Reviewed and Accepted

Interpretability and safety (workshop)

1. **P. M. Konrad**, T. Tanel, and S. Ayvaz (2026). *Self-Reports Do Not Identify Self-Models: An Identifiability Test for Counterfactual Reports*. Accepted (oral), *PhilML Workshop, ICML 2026*. [PDF]
2. **P. M. Konrad**, T. Tanel, and S. Ayvaz (2026). *Decoded but Unused: Instruction Tuning Routes Moral Framing into the Judgment Readout*. Accepted (poster), *ICML 2026 Workshop on Mechanistic Interpretability*. [PDF]
3. **P. M. Konrad**, T. Tanel, and S. Ayvaz (2026). *A Path Already Walked: On Inheriting Network-Neuroscience Tools for Mechanistic Interpretability*. Accepted (virtual poster), *ICML 2026 Workshop on Mechanistic Interpretability*. [PDF]

AI for software engineering (workshop)

4. **P. M. Konrad**, T. L. Adam, R. Terrenzi, and S. Ayvaz (2026). *Architecture Without Architects: How AI Coding Agents Shape Software Architecture*. Accepted, *SAGAI Workshop, IEEE ICSA 2026*. arXiv:2604.04990
 5. T. L. Adam, **P. M. Konrad**, R. Terrenzi, F. G. Lukas, R. Yilmaz, K. Sierszecki, and S. Ayvaz (2026). *CAKE: Cloud Architecture Knowledge Evaluation of Large Language Models*. Accepted, *KDA-AI Workshop, IEEE ICSA 2026*. arXiv:2604.05755
 6. R. Terrenzi, **P. M. Konrad**, T. L. Adam, and S. Ayvaz (2026). *A Reference Architecture for Agentic Hybrid Retrieval in Dataset Search*. Accepted, *SAML Workshop, IEEE ICSA 2026*. arXiv:2604.16394
- Conference proceedings*
7. N.-T. Thinh, Y. Du, **P. M. Konrad**, and A. Narechania (2026). *Fact-check Your Information (FYI): A Design Probe to Understand How People Actually Fact-check Data-Driven Articles*. Accepted, *IEEE VIS 2026*. [PDF]
 8. **P. M. Konrad**, T. Tanel, and S. Ayvaz (2025). *Beyond Major Floods: Deep Learning for Detecting Shallow Water Inundation in Agricultural Areas*. In *KES 2025*, *Procedia Computer Science*, 270, 301–310. [open access]

Preprints and Under Review

Interpretability and safety

9. **P. M. Konrad**, T. Tanyel, and S. Ayvaz (2026). *Routing Subspaces: Auditing Evaluation-to-Deployment Mismatch in Fine-Tuned Language Models*. Under review. [PDF]
10. **P. M. Konrad**, T. Tanyel, and S. Ayvaz (2026). *Acceptance Cards: A Four-Diagnostic Standard for Safe Fine-Tuning Defense Claims*. Under review. arXiv:2605.10575
11. **P. M. Konrad**, T. L. Adam, A. C. H. Merrild, R. de Rosa, R. Terrenzi, T. Tanyel, and S. Ayvaz (2026). *The Open-Box Fallacy: Why AI Deployment Needs a Calibrated Verification Regime*. Under review. arXiv:2605.10601
12. **P. M. Konrad**, Y. Du, and S. Ayvaz (2026). *Grounded Auditing as an Evaluation Policy: A Matched-Action Protocol and Stress Test*. Under review. [PDF]
13. **P. M. Konrad** and S. Ayvaz (2026). *When does chain-of-thought improve safety? Evidence from 18 models across 5 families*. Under review. [PDF]
14. **P. M. Konrad**, Y. Du, T. Tanyel, and S. Ayvaz (2026). *How Much of a Probe Direction Lives in the Input Embedding? An Audit of Four Published Claims*. Under review. [PDF]
15. **P. M. Konrad**, Y. Du, T. Tanyel, and S. Ayvaz (2026). *Abstention and Refusal Are Separable Directions in the Residual Stream*. Under review. [PDF]

Applied ML and other

16. **P. M. Konrad**, A.-A. Popa, Y. Sabzehmeidani, L. Zhong, M. Tripathy, A. Constantinescu, E. A. Liehn, and S. Ayvaz (2025). *Challenges in Deep Learning-Based Small Organ Segmentation: A Benchmarking Perspective for Medical Research with Limited Datasets*. Under revision, *Biomedical Signal Processing and Control*; arXiv:2509.05892.
17. **P. M. Konrad**, Y. Sabzehmeidani, A. Popa, and S. Ayvaz (2026). *Machine Learning in Gastrointestinal Tract Imaging: A Comprehensive Review of Techniques and Applications*. Preprint. Zenodo: 10.5281/zenodo.19426162
18. **P. M. Konrad**, C. H. Kunstmann-Olsen, J. Fiutowski, and S. Ayvaz (2025). *Non-Destructive Prediction of Fruit Ripeness and Firmness Using Hyperspectral Imaging and Lightweight Machine Learning Models*. Under review. arXiv:2604.22788

Research Experience

DataVISards, Hong Kong University of Science and Technology (HKUST) Jan 2026 – Present

Research Collaborator

Hong Kong SAR, China

- Collaborating with the DataVISards research group on machine learning and data visualisation projects.

Data and Intelligence Lab, SDU 

Sep 2024 – Present

Research Collaborator

Jan 2026 – Present

- Independently initiate and drive machine learning research projects: problem formulation, experimental design, pipeline development, manuscript writing, and venue selection.
- Co-author papers with Assoc. Prof. Serkan Ayvaz and lab members across remote sensing, medical imaging, NLP, and software systems.
- Lead the full research lifecycle end-to-end without external task assignment.

Research Assistant

Sep 2024 – Dec 2025

- Built end-to-end ML pipelines and contributed to study design, literature reviews, model development, and manuscript preparation.
- Contributed to grant proposal preparation and multimodal dataset management (remote sensing, hyperspectral, medical imaging).
- Co-authored the KES-2025 paper on deep learning for shallow water inundation detection.

Academic Service

Teaching Assistant, Artificial Intelligence (SDU) 

Jan 2026 – Jun 2026

- Teaching Assistant for the BSc-level course “Artificial Intelligence,” supporting exercise sessions and student supervision.
- Built an interactive AI course covering AI fundamentals with step-by-step animations, used as supplementary material for the course.
- Responsibilities include designing hands-on lab assignments, providing student support, and contributing to curriculum development.

Educational Committee, Software Engineering Programme, SDU

Student Member

Jun 2025 – Present
Sønderborg, Denmark

- Represent BSc Software Engineering students in curriculum and quality-assurance discussions, providing feedback on course content, assessment, and study environment.

Professional Experience

SaturoLabs

Founder

Jan 2026 – Present
Denmark

- Founded and currently leading SaturoLabs, building products at the intersection of AI and software engineering.
- Live products: claudeboyz.com, getproofz.com, dreambear.app.

Tutora ApS

CTO and Co-Founder

Jun 2024 – Nov 2024
Sønderborg, Denmark

- Led the end-to-end development of the company’s websites and core web application.
- Implemented the Shape Up product development framework to streamline technical execution.

Yeager GmbH

Co-CEO and Co-Founder

Nov 2022 – Jun 2024
Remote

- Co-founded the company and led development of *stabil.ai*, an AI-powered mobile app for personalized powerlifting training.
- Engineered algorithms to personalize training plans using MRV/MEV and dynamic real-time feedback.
- Oversaw the full product lifecycle, from UX/UI concept and design to full-stack implementation.

Transition Year: Independent Study & Prototyping

Independent

Oct 2021 – Oct 2022
Germany/Denmark

- Completed intensive coding bootcamps and self-directed study in software engineering to transition from military service into the technology field.
- Began prototyping early concepts for a training-app product, which later evolved into the product developed at Yeager GmbH.

Bundeswehr (German Armed Forces)

Staff Duty Soldier

Oct 2017 – Sep 2021
Glücksburg, Germany

- Led a small HR team responsible for the administration of over 600 soldiers.
- Streamlined administrative processes and document workflows in a high-stakes naval command environment.

Telecommunications Company

Customer Service Consultant

Jul 2016 – Sep 2017
Germany

- Provided front-line customer support for mobile and broadband services.

Skills

Programming Languages: Python, JavaScript, SQL

ML Libraries: PyTorch, Hugging Face, TransformerLens, Scikit-learn, Pandas, NumPy

Data Visualization: Matplotlib, Seaborn

ML Engineering: Data & Feature Engineering, Hyperparameter Tuning (Optuna, Weights & Biases)

Sweeps), Experiment Tracking (Weights & Biases)

Web Development: React, React Native, Node.js, Express, HTML/CSS, Streamlit

Tools & Platforms: Git, Docker, Google Cloud Platform, Vercel

Methodologies: Scrum, Shape Up

Languages

English (Professional working)

German (Native or bilingual)

Thai (Native or bilingual)

Danish (Elementary)

Honors and Awards

William Demant Fonden Grant May 2026

Received a grant from the William Demant Foundation supporting graduate studies.

Top-10 Placement, Danish National Championship in AI (DM i AI) 🌐 2025

National AI competition, individual competitor

- Secured a top-10 placement while competing solo against more than 50 teams in Denmark's largest AI competition, organized by the Danish Society for Artificial Intelligence.

Top-10 Placement, Danish National Championship in AI (DM i AI) 🌐 2024

National AI competition, individual competitor

- Achieved top-tier placement in a national AI competition for students, organized by the Danish Society for Artificial Intelligence.

1st Place, SDU Case Competition, SDU Sønderborg 🌐 🌐 2023

University case competition on sustainability

- Awarded 1st place in a 48-hour competition focused on sustainability with real-world cases from companies such as Danfoss and Linak.

Formal Recognition for Exemplary Service, Bundeswehr (German Armed Forces) 2021

Military commendation

- Formally commended for setting a benchmark in dedication and professionalism for enlisted personnel.

Performance Bonus for Outstanding Achievement, Bundeswehr (German Armed Forces) 2021

Military achievement bonus

- Awarded a financial bonus for sustained, excellent performance and independent handling of complex tasks.

Certifications

Venture Capital Explorer Programme, Accelerace & BII 2026

Intensive 4-day program covering VC investment process, founder sourcing, startup assessment, venture risk evaluation, investment committee operations, exits, and case-study simulations.

Machine Learning, Stanford Online 2024

Foundational C# with Microsoft 2024

Google UX Design Professional, Google 2023

React Native Course 2023

Full-Stack Engineer Career Path, Codecademy 2023

Certified Specialist for Real Estate Loan Brokerage (IHK) 2022