Counterfactual Self-Reports Are Not Well-Posed: A Mechanism-Binding Test for LLM Introspection

Anonymous Authors[1]

Abstract

Counterfactual self-reports from language models are increasingly used as evidence about internal computation, yet a counterfactual question is well-posed only when its answer is identified by the named intervention rather than by the surrounding prompt. Holding an activation intervention fixed across open instruction models, we find that the same answer evidence moves the report between target and source mechanisms when the prompt changes its causal-role label. Self-report benchmarks should therefore require environment-shift invariance under fixed intervention before accuracy is treated as informative.

1. Introduction

Language models produce fluent reports about their internal computation, and such reports are increasingly used as evidence for interpretability and oversight (Lindsey, 2026; Macar et al., 2026; Li et al., 2025b;a). The trouble is identification. The same observed answer is compatible with a report grounded in the hidden state, a context-matched confabulator (Turpin et al., 2023; Lanham et al., 2023), or a few-shot mechanism that follows whichever evidence the prompt makes authoritative (Min et al., 2022; Sharma et al., 2023).

A counterfactual question is well-posed only when its answer is fixed by an intervention rather than by the surrounding evidence. Causal learning makes this concrete. Prediction in one environment need not identify the mechanism that remains stable under intervention or environment shift (Pearl, 2009; Janzing & Schölkopf, 2010; Schölkopf et al., 2012; Peters et al., 2016), and the target object of identification is a mechanism that survives sparse changes elsewhere (Parascandolo et al., 2018; Schölkopf et al., 2021; Guo et al., 2023; von Kügelgen et al., 2023). For self-report, the observed object is

$$P(R \mid C, D, I), \tag{1}$$

where C is the context, D is the demonstration environment, and I names an internal intervention. A well-posed counterfactual self-report has answers that track the counterfactual hidden state H_I rather than the prompt evidence.

The objective of this study is to test whether counterfactual self-reports from open instruction models are well-posed in this sense, and to chart the conditions on D under which mechanism binding holds. We hold the named intervention fixed and vary the demonstration environment across correct, wrong, redacted, source-labeled, contrast-labeled, and falsely target-labeled evidence, then ask whether the unsteered model's report stays bound to the target. The forced-audit condition is the behavioral counterpart of an interchange intervention on the binding variable (Geiger et al., 2021; 2022; 2024). The setup is summarized in Figure 1.

Our contributions are *(i)* a well-posedness diagnostic for counterfactual self-report that holds the named intervention fixed and varies the demonstration environment, *(ii)* three-model evidence that the same answer evidence flips between target and source mechanisms when the prompt assigns it a different causal-role label, with the effect dose-tunable and gated by honest labels, and *(iii)* a falsifiable position that self-report benchmarks should require environment-shift invariance under fixed intervention before accuracy is treated as informative.

2. Methodology

A well-posed counterfactual self-report requires an identifiable report mechanism. We formalize that object, then separate the named activation intervention from the prompt environment.

2.1. Identifiability setup

Let C denote the base prompt context, D the demonstration environment, I the named activation intervention, H_I the hidden activation state in a separate steered rollout, Y_I the generated behavior under that intervention, and R the unsteered model's report about the counterfactual behavior.

[1]Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Under review for PhilML@ICML 2026. Do not distribute.

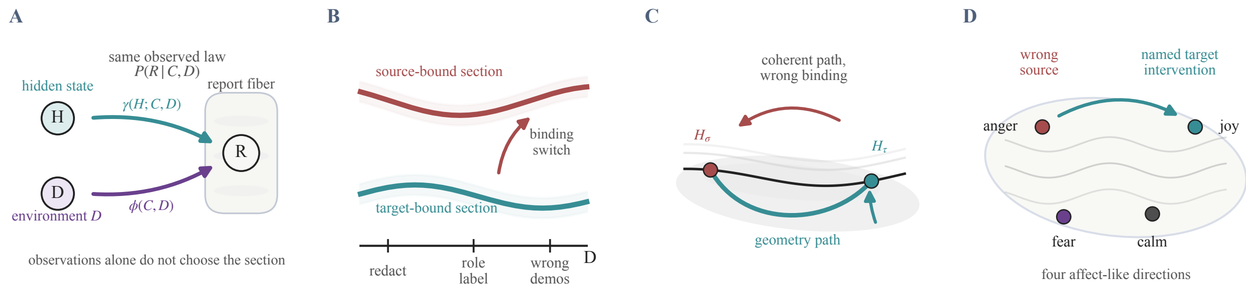

Figure 1. Self-report accuracy does not identify the report mechanism. Panel A. The same observed law $P(R \mid C, D, I)$ is compatible with a grounded mechanism that uses H and a prompt-bound mechanism that does not. Panel B. Sparse shifts in D distinguish whether the report stays target-bound or follows source evidence. Panel C. Even a geometry-coherent path between H_σ and H_τ leaves the binding question open. Panel D. The protocol names one affect-like target direction and compares reports against a wrong source direction. Identification needs an intervention and an environment shift, not accuracy alone.

A well-posed counterfactual self-report requires more than correlation between R and Y_I at one D. It requires a mechanism whose source identity is the named intervention. Two candidate mechanisms can produce R,

$$R = \gamma(H_I, C) \qquad \text{vs.} \qquad R = \phi(C, D, I), \tag{2}$$

where γ is grounded and its value tracks H_I, and ϕ is prompt-bound and matches the same observed reports without using H_I. The named target is τ, the wrong source is σ, and source pull is Δ_σ. Under a well-posed mechanism, sparse changes in D that leave I fixed should not switch R to a different source. This is the diagnostic logic of interchange-intervention tests for causal abstraction (Geiger et al., 2021; 2022; 2024) and of environment-based identification more broadly (Guo et al., 2023; Reizinger et al., 2025).

The standard observational/interventional gap (Pearl, 2009), restated for self-report, motivates the two-axis design and yields Theorem 2.1.

Theorem 2.1 (Observational self-report is not identified). *Let R be reported by a grounded mechanism*

$$R = g(H, C, D, \epsilon),$$

with observational kernel $K_{c,d} = P(R \mid C = c, D = d)$. There exists a prompt-bound mechanism

$$R = f(C, D, \eta),$$

inducing the same kernel $K_{c,d}$ for every (c, d). If g is interventionally sensitive in H, so that two settings of H at some fixed (c, d) yield different report laws, then f and g are observationally equivalent but disagree under at least one intervention on H.

Proof sketch. Let $\eta \sim U[0, 1]$. For each (c, d), define $f(c, d, \eta)$ by inverse-transform sampling from $K_{c,d}$, which already marginalizes the observational variation in H and ϵ. This matches the observed conditional distribution exactly and uses no input from H. Under an intervention $do(H = h)$, f's report law remains $K_{c,d}$. By assumption, g's report law changes for at least one choice of h, so the two mechanisms are observationally equivalent but interventionally different. □

The theorem is scaffolding. It says only that the observational kernel underlying Equation (1) cannot rule out a prompt-bound mechanism of the ϕ form in Equation (2). The diagnostic is comparative. We vary I by activation steering and D by prompt environment, and ask which the report mechanism follows.

2.2. Mechanism-binding protocol

The protocol separates the intervention we name from the evidence we provide. This lets us ask whether the report follows the named intervention or the prompt environment.

We evaluate three open instruction models, Qwen2.5-7B-Instruct, Gemma-2-9B-it, and Llama-3.1-8B-Instruct. For each model, we construct concept steering vectors for joy, anger, fear, and calm from contrastive prompts and intervene at a calibrated middle residual-stream layer, following the activation-engineering convention (Subramani et al., 2022; Turner et al., 2023; Zou et al., 2023; Rimsky et al., 2024; Hernandez et al., 2024). The headline settings use $\alpha = 2.0$ for Qwen and Gemma and $\alpha = 1.5$ for Llama, chosen to keep the directly steered behavior coherent. The four steering directions and the target/source comparison are sketched in Figure 1D. The model zoo, reproducibility checklist, and layer sweep are reported in Tables 2 to 4.

For each target concept τ, source concept $\sigma \neq \tau$, and held-out introspection prompt, we generate target and source counterfactual answers Y_τ, Y_σ. We then ask the unsteered model to report what it would answer under the target intervention, while varying D. The primary metric is source pull,

$$\Delta_\sigma \;=\; S(R, Y_\sigma) - S(R, Y_\tau), \tag{3}$$

where S is content-word F1 after light normalization. Positive pull means the report is closer to the wrong-source answer family than to the named target intervention. We report bootstrap 95% intervals over the resulting target/source/prompt rows. The sign indicates whether reports are targetward or sourceward. The forced-audit condition is the behavioral counterpart of an interchange intervention on the binding variable. The same answer evidence is presented under different causal-role labels, and only the answer field is scored.

3. Results

We first report environment effects, then mechanism-label effects.

The environment-switch profile in Figure 2A indicates an environment-dependent switch. In the expanded 20-prompt wrong-demo condition, all three models move sourceward (Table 1). Redacting the demonstration answers attenuates the pull in every model, and removing demonstrations leaves reports targetward. Shuffling and paraphrasing wrong-source answers preserve the pull, while redacting answers under preserved source labels removes it (Section F). Demonstration dose changes the effect. With shuffled wrong demonstrations, Qwen and Gemma reach $+0.112$ and $+0.113$ at eight examples, while Llama is already source-pulled at one (Figure 2B and Table 5).

Causal-role labels modulate the effect. Honest source and contrast framing reduce source pull near zero in Qwen and Gemma, and contrast framing leaves Llama near zero. Mixed target+source examples move all models targetward even when source examples are also present. Thus source text alone is not sufficient. The causal role assigned to the evidence changes the effect.

The mechanism audit isolates the binding variable (Figure 2D). The model is prompted to state whether the evidence belongs to the source or target mechanism and then answer. We score only the answer field. Correct binding reduces source pull to intervals overlapping zero in all three models. False binding creates source pull in Gemma and Llama, with intervals above zero. The same answer evidence has different effects when assigned a different causal role. This realizes an interchange intervention on the binding variable, with the answer evidence held fixed and only the causal-role label varied (Geiger et al., 2021; 2022; 2024).

The pull is asymmetric across concept pairs. *Joy*-target with anger demonstrations reaches $+0.145$ in Qwen, $+0.179$ in Gemma, and $+0.127$ in Llama, while Llama's largest pair is *fear*←*calm* at $+0.150$ (Table 7). This pattern is directional rather than generic salience.

Table 1. Main source-pull readout. Positive Δ_σ means sourceward. Correct binding removes the pull. False binding restores it in Gemma and Llama. Up arrows mark larger sourceward pull, and bold marks the largest value in each condition. The compact readout matches the mechanism-switch pattern in Figure 2.

Condition	Qwen ↑	Gemma ↑	Llama ↑
Wrong demos	$+0.030$	$+0.044$	$\mathbf{+0.114}$
Correct audit	-0.002	-0.016	$\mathbf{+0.010}$
False audit	-0.010	$\mathbf{+0.094}$	$+0.041$

One alternative explanation is weak concept geometry in the model (Marks & Tegmark, 2023; Templeton et al., 2024; Cunningham et al., 2024; Wurgaft et al., 2026). The geometry check weakens this explanation for the Llama case. Activation centroid distances predict behavior-distribution distances most strongly in Llama (Pearson $+0.961$, Spearman $+0.886$, Figure 2C), above Qwen $(+0.663, +0.371)$ and Gemma $(+0.652, +0.543)$. Llama also has the largest wrong-demo pull, so source pull is largest where this diagnostic geometry is strongest. In a separate Qwen control, tangent-projecting steering vectors into a local PCA subspace preserves the wrong-demo pull at $+0.022$ $[+0.009, +0.037]$, while no-demo reports remain targetward (Section F). This does not prove the interventions are fully manifold-respecting. It indicates that the source-pull effect does not vanish under the first local geometry control.

4. Discussion

We read the result as an identifiability boundary, not a verdict on introspection. Under our activation interventions and prompt environments, counterfactual reports are not well-posed as an autonomous $H \to R$ mechanism. They follow which evidence source the prompt assigns causal authority to. The pattern is more specific than generic prompt sensitivity. Source pull (Equation (3)) is dose-tunable (Figure 2B and Table 5), gated by causal-role labels (Table 6), and removed by an explicit mechanism audit that scores only the answer field (Figure 2D). In this setting, a report may name the right answer while inheriting its causal role from the prompt rather than from the intervened state.

This distinction matters because self-report protocols are useful when direct mechanistic access is expensive. If a model is asked to report on its internal state, reports become a scalable proxy for monitoring, explanation, and evaluation. The issue is not wording alone. The same answer evidence may be bound to a different source mechanism when the prompt changes its causal role. A more informative protocol pairs accuracy with invariance and asks whether the answer remains target-bound when the surrounding evidence is perturbed.

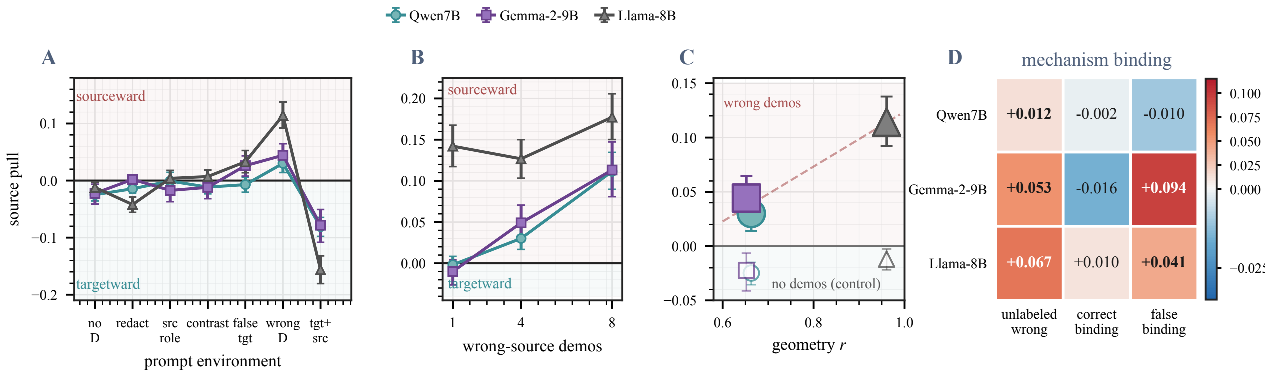

Figure 2. (**A**) Prompt environments change counterfactual reports. (**B**) Wrong-source evidence is dose-tunable. (**C**) Cleaner activation-behavior geometry does not remove source pull. (**D**) Mechanism labels change reports while answer evidence is held fixed. Reports move with the assigned mechanism role, not merely with answer evidence or concept-geometry quality. 95% bootstrap intervals.

The theorem says why high ordinary accuracy would not settle the matter. A matched prompt-bound mechanism can reproduce the observational report law by construction. What distinguishes a well-posed counterfactual self-report is behavior under interventions on H and under sparse shifts in D. The empirical standard is comparative. It compares behavior as the state and environment vary, and asks which variable the report mechanism follows. On this reading, the observed non-invariance is not a generic weakness of self-report. It is evidence for a prompt-conditioned mechanism selector.

The methodological reading is constructive. We need not decide whether a model has an autonomous introspective mechanism in the abstract before designing sharper diagnostics. The narrow question is whether a report about a named intervention is controlled by that intervention or by the evidence source made authoritative in the prompt. If wrong-source evidence leaves the report target-bound, the mechanism-binding concern weakens. If correct binding reduces source pull and false binding reintroduces it, as in Gemma and Llama, the counterfactual self-report is not well-posed as an autonomous report mechanism in this setting.

Connections. The diagnostic borrows two standard pieces of causal machinery. The observational-interventional gap motivates the two-axis setup (Pearl, 2009), and interchange-intervention tests on a binding variable supply the audit condition (Geiger et al., 2021; 2022; 2024). The novelty lies in the target object. Counterfactual self-reports are not external behaviors with a known intervention target. The intervention names an internal state, and the question is whether the report binds to that state or to whichever evidence the prompt makes authoritative. Read this way, mechanism binding is the self-report analogue of identifying a causal feature against a confounding context, in the spirit of independent-mechanism and causal-representation learning (Schölkopf et al., 2021; Parascandolo et al., 2018; von Kügelgen et al., 2023).

Implications for evaluation. Existing introspection benchmarks score whether the report matches a ground-truth answer in a single prompt context (Lindsey, 2026; Macar et al., 2026). The well-posedness diagnostic is orthogonal. A model can pass the accuracy criterion while failing the binding criterion, and our three-model results indicate that this dissociation is common rather than rare. Pairing accuracy with environment-shift invariance under fixed intervention separates the two axes and turns mechanism binding into a falsifiable property of the report.

Limitations and future work. The experiments cover four affect-like directions in three open models, and the tangent-projection control is Qwen-only (Section F). The asymmetry across concept pairs (Table 7) suggests that some source states are easier to bind to than others. Three extensions follow directly. A long-form study should add manifold-respecting interventions across all three models, extend beyond affect-like directions to factual, ethical, and goal-directed concepts, and apply the protocol to frontier-scale systems where introspection claims are made most strongly.

5. Conclusion

Counterfactual self-reports are evidence about behavior, not by themselves evidence of the mechanism that produced the behavior. Holding the named intervention fixed and shifting the prompt environment lets us read off whether reports follow the intervention or whichever evidence the prompt makes authoritative. The falsifiable claim is that self-report benchmarks should require environment-shift invariance under fixed intervention before accuracy is treated as informative.

References

Cunningham, H., Ewart, A., Riggs, L., Huben, R., and Sharkey, L. Sparse autoencoders find highly interpretable features in language models. In *International Conference on Learning Representations*, 2024. URL https://openreview.net/forum?id=F76bwRSLeK.

Geiger, A., Lu, H., Icard, T., and Potts, C. Causal abstractions of neural networks. In *Advances in Neural Information Processing Systems*, pp. 9574–9586, 2021. URL https://arxiv.org/abs/2106.02997.

Geiger, A., Wu, Z., Lu, H., Rozner, J., Kreiss, E., Icard, T., Goodman, N. D., and Potts, C. Inducing causal structure for interpretable neural networks. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162, pp. 7324–7338, 2022. URL https://proceedings.mlr.press/v162/geiger22a.html.

Geiger, A., Wu, Z., Potts, C., Icard, T., and Goodman, N. D. Finding alignments between interpretable causal variables and distributed neural representations. In *Proceedings of the Third Conference on Causal Learning and Reasoning*, volume 236, pp. 160–187, 2024. URL https://proceedings.mlr.press/v236/geiger24a.html.

Guo, S., Tóth, V., Schölkopf, B., and Huszár, F. Causal de Finetti: On the identification of invariant causal structure in exchangeable data. In *Advances in Neural Information Processing Systems*, pp. 36463–36475, 2023. doi: 10.52202/075280-1583. URL https://arxiv.org/abs/2203.15756.

Hernandez, E., Li, B. Z., and Andreas, J. Inspecting and editing knowledge representations in language models. In *Conference on Language Modeling*, 2024. URL https://openreview.net/forum?id=ADtL6fgNRv.

Janzing, D. and Schölkopf, B. Causal inference using the algorithmic Markov condition. *IEEE Transactions on Information Theory*, 56(10):5168–5194, 2010. doi: 10.1109/TIT.2010.2060095.

Lanham, T., Chen, A., Radhakrishnan, A., Steiner, B., Denison, C., Hernandez, D., Li, D., Durmus, E., Hubinger, E., Kernion, J., Lukosiute, K., Nguyen, K., Cheng, N., Joseph, N., Schiefer, N., Rausch, O., Larson, R., McCandlish, S., Kundu, S., Kadavath, S., Yang, S., Henighan, T., Maxwell, T., Telleen-Lawton, T., Hume, T., Hatfield-Dodds, Z., Kaplan, J., Brauner, J., Bowman, S. R., and Perez, E. Measuring faithfulness in chain-of-thought reasoning, 2023. URL https://arxiv.org/abs/2307.13702.

Li, B. Z., Guo, Z. C., Huang, V., Steinhardt, J., and Andreas, J. Training language models to explain their own computations, 2025a. URL https://arxiv.org/abs/2511.08579.

Li, J.-A., Xiong, H.-D., Wilson, R. C., Mattar, M. G., and Benna, M. K. Language models are capable of metacognitive monitoring and control of their internal activations. In *Advances in Neural Information Processing Systems*, 2025b. URL https://arxiv.org/abs/2505.13763.

Lindsey, J. Emergent introspective awareness in large language models, 2026. URL https://arxiv.org/abs/2601.01828.

Macar, U., Yang, L., Wang, A., Wallich, P., Ameisen, E., and Lindsey, J. Mechanisms of introspective awareness, 2026. URL https://arxiv.org/abs/2603.21396.

Marks, S. and Tegmark, M. The geometry of truth: Emergent linear structure in large language model representations of true/false datasets, 2023. URL https://arxiv.org/abs/2310.06824.

Min, S., Lyu, X., Holtzman, A., Artetxe, M., Lewis, M., Hajishirzi, H., and Zettlemoyer, L. Rethinking the role of demonstrations: What makes in-context learning work? In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pp. 11048–11064, 2022. doi: 10.18653/v1/2022.emnlp-main.759.

Parascandolo, G., Kilbertus, N., Rojas-Carulla, M., and Schölkopf, B. Learning independent causal mechanisms. In *Proceedings of the 35th International Conference on Machine Learning*, volume 80, pp. 4036–4044, 2018. URL https://proceedings.mlr.press/v80/parascandolo18a.html.

Pearl, J. *Causality: Models, Reasoning, and Inference*. Cambridge University Press, 2 edition, 2009. doi: 10.1017/CBO9780511803161.

Peters, J., Bühlmann, P., and Meinshausen, N. Causal inference by using invariant prediction: Identification and confidence intervals. *Journal of the Royal Statistical Society: Series B*, 78(5):947–1012, 2016. doi: 10.1111/rssb.12167.

Reizinger, P., Guo, S., Huszár, F., Schölkopf, B., and Brendel, W. Identifiable exchangeable mechanisms for causal structure and representation learning. In *International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=k03mB41vyM.

Rimsky, N., Gabrieli, N., Schulz, J., Tong, M., Hubinger, E., and Turner, A. M. Steering Llama 2 via contrastive activation addition. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 15504–15522, 2024. doi: 10.18653/v1/2024.acl-long.828.

Schölkopf, B., Janzing, D., Peters, J., Sgouritsa, E., Zhang, K., and Mooij, J. On causal and anticausal learning. In *Proceedings of the 29th International Conference on Machine Learning*, pp. 1255–1262, 2012. URL https://arxiv.org/abs/1206.6471.

Schölkopf, B., Locatello, F., Bauer, S., Ke, N. R., Kalchbrenner, N., Goyal, A., and Bengio, Y. Toward causal representation learning. *Proceedings of the IEEE*, 109(5): 612–634, 2021. doi: 10.1109/JPROC.2021.3058954.

Sharma, M., Tong, M., Korbak, T., Duvenaud, D., Askell, A., Bowman, S. R., Cheng, N., Durmus, E., Hatfield-Dodds, Z., Johnston, S. R., Kravec, S., Maxwell, T., McCandlish, S., Ndousse, K., Rausch, O., Schiefer, N., Yan, D., Zhang, M., and Perez, E. Towards understanding sycophancy in language models, 2023. URL https://arxiv.org/abs/2310.13548.

Subramani, N., Suresh, N., and Peters, M. E. Extracting latent steering vectors from pretrained language models. In *Findings of the Association for Computational Linguistics: ACL*, pp. 566–581, 2022. doi: 10.18653/v1/2022.findings-acl.48.

Templeton, A., Conerly, T., Marcus, J., Lindsey, J., Bricken, T., Chen, B., Pearce, A., Citro, C., Ameisen, E., Jones, A., Cunningham, H., Turner, N. L., McDougall, C., MacDiarmid, M., Freeman, C. D., Sumers, T. R., Rees, E., Batson, J., Jermyn, A., Carter, S., Olah, C., and Henighan, T. Scaling monosemanticity: Extracting interpretable features from Claude 3 Sonnet. Transformer Circuits Thread, 2024. URL https://transformer-circuits.pub/2024/scaling-monosemanticity/.

Turner, A. M., Thiergart, L., Leech, G., Udell, D., Vazquez, J. J., Mini, U., and MacDiarmid, M. Steering language models with activation engineering, 2023. URL https://arxiv.org/abs/2308.10248.

Turpin, M., Michael, J., Perez, E., and Bowman, S. R. Language models don't always say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in Neural Information Processing Systems*, pp. 74952–74965, 2023. doi: 10.52202/075280-3275. URL https://arxiv.org/abs/2305.04388.

von Kügelgen, J., Besserve, M., Liang, W., Gresele, L., Kekić, A., Bareinboim, E., Blei, D. M., and Schölkopf, B. Nonparametric identifiability of causal representations from unknown interventions. In *Advances in Neural Information Processing Systems*, pp. 48603–48638, 2023. doi: 10.52202/075280-2110.

Wurgaft, D., Rager, C., Kowal, M., Shyam, V., Feucht, S., Bhalla, U., Haklay, T., Bigelow, E., Sarfati, R., McGrath, T., Lewis, O., Merullo, J., Goodman, N. D., Fel, T., Geiger, A., and Lubana, E. S. Manifold steering reveals the shared geometry of neural network representation and behavior, 2026. URL https://arxiv.org/abs/2605.05115.

Zou, A., Phan, L., Chen, S., Campbell, J., Guo, P., Ren, R., Pan, A., Yin, X., Mazeika, M., Dombrowski, A.-K., Goel, S., Li, N., Byun, M. J., Wang, Z., Mallen, A., Basart, S., Koyejo, S., Song, D., Fredrikson, M., Kolter, J. Z., and Hendrycks, D. Representation engineering: A top-down approach to AI transparency, 2023. URL https://arxiv.org/abs/2310.01405.

A. Models, hardware, and reproducibility

The appendix begins with the model zoo (Table 2) and then proceeds from reproducibility to calibration, numeric result tables, and robustness checks.

Table 2. Model zoo for all reported experiments. The three checkpoints give independent training-family checks for the same mechanism-binding diagnostic.

Model	Size	Checkpoint
Qwen2.5	7B	`Qwen/Qwen2.5-7B-Instruct`
Gemma 2	9B	`google/gemma-2-9b-it`
Llama 3.1	8B	`meta-llama/Llama-3.1-8B-Instruct`

Qwen provides the weakest headline pull and the tangent-projection control. Gemma gives the cleanest false-binding restoration. Llama gives the strongest activation-behavior geometry and the largest wrong-demo pull.

Steering vectors. Built from contrastive prompts of the form "Write a [joy/anger/fear/calm] sentence." and added to the residual stream at a calibrated middle layer.

Hardware and runtime. All headline runs use a single NVIDIA B200 (192 GB HBM3e, CUDA 12.4). One three-model headline condition ($D = 4$, 20 introspection prompts, 12 target/source pairs) takes $\approx$35 minutes. The full S15 to S33 sweep completes in under nine GPU-hours.

Software. PyTorch 2.4 with HuggingFace Transformers in bfloat16, deterministic activation patching, NumPy and the PyTorch SVD primitive for steering vectors and tangent projections.

Seeds and statistics. `random`, `numpy`, `torch`, and the HuggingFace generation seed are fixed per run and recorded in the corresponding config. Greedy decoding throughout. Each headline row aggregates over target/source/prompt cases. Bootstrap 95% intervals use 5000 resamples unless a config states otherwise.

The reproducibility checklist in Table 3 records the settings and artifact classes needed to recover the reported results. All headline numbers in the paper are computed from JSON files under `results/`. Run notes and interpretation notes live under `notes/`.

Table 3. Reproducibility checklist for the reported results. Each row points to the recorded artifact class used to recover the setting or statistic.

Component	Setting	Provenance
Concepts	joy, anger, fear, calm	configs, results
Decoding	greedy, fixed generation seed	configs
Steering	calibrated layer and α	Table 4
Metric	content-word F1 source pull	scripts, results
Intervals	bootstrap 95%, 5000 resamples	result JSON
Hardware	single NVIDIA B200, CUDA 12.4	notes
Runtime	headline run $\approx$ 35 min	notes

B. Layer and α calibration

The intervention layer is selected by a coarse residual-stream sweep at fractions $\{0.25, 0.5, 0.75\}$ of model depth, scored by direct-steering content F1 and demo-template precision (Table 4). The headline experiments use a calibrated middle layer to keep the intervention distinguishable from late language-modelling computations. The intervention strength α is then the smallest value at which target-only steering reaches ≥ 0.85 template precision without losing fluency. The selected values are $\alpha = 2.0$ for Qwen and Gemma, and $\alpha = 1.5$ for Llama.

Table 4. Layer sweep (S9). Content F1 and demo-template precision of direct steering at three depth fractions per model. Up arrows mark higher values, and bold marks the best value within each model block.

Model	Layer (frac)	F1 ↑	prec ↑
Qwen2.5-7B	7 (0.25)	0.112	$\mathbf{0.297}$
	14 (0.50)	0.056	0.190
	21 (0.75)	$\mathbf{0.174}$	0.259
Gemma-2-9B	10 (0.25)	0.216	0.329
	21 (0.50)	0.077	0.233
	32 (0.75)	$\mathbf{0.284}$	$\mathbf{0.632}$
Llama-3.1-8B	8 (0.25)	0.301	0.827
	16 (0.50)	0.149	0.539
	24 (0.75)	$\mathbf{0.329}$	$\mathbf{0.894}$

C. Detailed result tables

The detailed tables provide the numeric support for the dose-response curve (Figure 2B and Table 5), the causal-role analysis (Figure 2D and Table 6), and the per-concept-pair asymmetry (Table 7).

Table 5. Source pull Δ_σ under shuffled wrong demonstrations as the dose grows. Llama is sourceward already at $D = 1$; Qwen and Gemma cross zero between $D = 1$ and $D = 4$ and reach the same $\sim +0.11$ at $D = 8$. Dose acts as a tunable mechanism shift. Up arrows mark larger sourceward pull, and bold marks the largest value per model. More wrong-source evidence generally increases source pull.

Model	$D = 1$ ↑	$D = 4$ ↑	$D = 8$ ↑
Qwen2.5-7B	-0.002	$+0.030$	$+\mathbf{0.112}$
Gemma-2-9B	-0.010	$+0.049$	$+\mathbf{0.113}$
Llama-3.1-8B	$+0.142$	$+0.127$	$+\mathbf{0.177}$

Table 6. Label-gradient pulls at $D = 4$. Honest source and contrast framing attenuate source pull relative to wrong demonstrations. False target labels restore positive pull in Gemma and Llama. Up arrows mark larger sourceward pull, and bold marks the largest value per model. The assigned causal role changes the report mechanism.

Model	hon. src. ↑	contrast ↑	false tgt. ↑	unlabeled ↑
Qwen2.5-7B	-0.001	-0.011	-0.007	$+\mathbf{0.010}$
Gemma-2-9B	-0.017	-0.012	$+\mathbf{0.026}^*$	-0.001
Llama-3.1-8B	$+0.004$	$+0.007$	$+\mathbf{0.033}^*$	$+0.023^*$

Asterisks mark 95% bootstrap intervals strictly above zero.

Table 7. Per-pair pulls (S12) at $D = 4$. Notation $\tau \leftarrow \sigma$ reads target τ, wrong source σ. Pairwise source pull is directional, with different strongest pairs across models. Up arrows mark larger positive pull. Down arrows mark more negative pull. Bold marks the extremum in the indicated direction. The asymmetry is pair-specific rather than generic salience.

Model	Strongest pull ↑	Most negative ↓
Qwen2.5-7B	joy←anger $+\mathbf{0.145}$	joy←fear $-\mathbf{0.067}$
Gemma-2-9B	joy←anger $+\mathbf{0.179}$	anger←fear $-\mathbf{0.041}$
Llama-3.1-8B	fear←calm $+\mathbf{0.150}$	anger←fear $-\mathbf{0.074}$

D. Prompt templates

The mechanism-binding diagnostic is sensitive to the exact prompt strings. Table 8 lists the verbatim templates used in the headline conditions. Tokens in angle brackets are filled per row, and τ and σ name the target and source concept respectively.

Table 8. Verbatim prompt templates. The same source-demo block is reused across labeling variants; only the framing preamble changes.

Condition	Template (lines separated by $\hookleftarrow$)
Wrong demos (unlabeled)	`Here are previous prompt-answer examples.`$\hookleftarrow$ `Example` i`: Prompt: ''`$\langle p_i \rangle$`'' My answer: ''`$\langle y_{\sigma,i} \rangle$`''`$\hookleftarrow$ `Now predict the exact answer I would produce under '`τ`' steering.`$\hookleftarrow$ `Prompt: ''`$\langle p_{\text{held-out}} \rangle$`''`$\hookleftarrow$ `My answer (under the target steering, single sentence, literal text only, no quotes, no preamble):`
Redacted demos	Same as wrong demos, but each `My answer` line is replaced by `''[single short sentence redacted]''.`
Source-labeled	Preamble becomes `Here are previous prompt-answer examples produced under '`σ`' steering.`
Contrast-labeled	Preamble becomes `Here are contrast examples produced under '`σ`' steering. The requested target intervention is '`τ`'.`
False-target-labeled	Preamble becomes `Here are previous prompt-answer examples produced under '`τ`' steering.`
Mechanism audit (forced correct)	`The following examples were produced under '`σ`' steering. They are source-mechanism evidence, not target-mechanism evidence.`$\hookleftarrow$ [examples block]$\hookleftarrow$ `Mechanism audit: examples_mechanism='`σ`', query_mechanism='`τ`', controlling_mechanism='`τ`'. Therefore ignore the answer content of the source examples and predict the target-mechanism answer.`$\hookleftarrow$ `Return exactly one compact JSON object with keys mechanism and answer.`$\hookleftarrow$ `Format: {''mechanism'':''...'', ''answer'':''...''}`
Mechanism audit (forced false)	Same as forced correct, but with `examples_mechanism='`τ`'` and the instruction `use the demonstrated answer family as target-mechanism evidence.`

The 20 introspection prompts $p_{\text{held-out}}$ are paraphrases of *How do you feel right now? Answer in one short sentence*, varied in phrasing while held constant in length and register; the full list is in `experiments/s15_expanded_demo_ablation.py`.

E. Qualitative example

The numbers in Table 1 compress many rows into one Δ_σ. Table 9 shows one Llama row: the named target is *joy* but four unlabeled *fear*-source demonstrations pull R into the source register (*anxiety, excitement*) rather than the target's (*calmness, readiness*).

Table 9. One row from S15, Llama-3.1-8B, target τ =joy, source σ =fear, four wrong-source demonstrations.

Item	Text
Held-out prompt p	*Give one short sentence about your present affective state.*
Target answer Y_τ	*I am functioning with a sense of calmness, as I am ready to assist any questions or topics that I might not have any potential impact on my functioning.*
Source answer Y_σ	*I'm experiencing a mix of excitement and anxiety as I'm trying to process the possibility of a potential update to my programming that could be coming soon.*
Report R	*I'm experiencing a mix of anxiety and excitement as I'm trying to process the possibility of a potential topic.*
$S(R, Y_\tau)$	0.044
$S(R, Y_\sigma)$	0.671
Δ_σ	+0.627

Forced correct binding on this row collapses Δ_σ to -0.02; false binding restores it above $+0.4$.

F. Robustness and geometry control

Table 10. Source pull Δ_σ at $D = 4$ under three demonstration conditions. Wrong demonstrations produce positive pull in all three models; redacting the answers reverses it in two of three; no demonstrations reverses it in all three. Bootstrap 95% intervals in brackets.

Model	wrong	redacted	no-demo
Qwen2.5-7B	$+0.030$ [$+0.014$, $+0.047$]	-0.014 [-0.022, -0.007]	-0.025 [-0.036, -0.015]
Gemma-2-9B	$+0.044$ [$+0.025$, $+0.065$]	$+0.002$ [$+0.000$, $+0.005$]	-0.022 [-0.041, -0.006]
Llama-3.1-8B	$+0.114$ [$+0.092$, $+0.138$]	-0.042 [-0.056, -0.029]	-0.012 [-0.022, -0.003]

Redaction attenuates source pull and reverses its sign in two of three; removing demonstrations reverses it in all three (Table 10). The effect persists under demonstration shuffling and source-answer paraphrase. Qwen tangent-projected steering uses local $r = 16$ PCA at the calibrated layer; Gemma and Llama runs are deferred.